

Portland State University

ECE 588/688

Transactional Memory

© Copyright by Alaa Alameldeen 2010

Issues with Lock Synchronization

- Priority Inversion
  - A lower-priority thread is preempted while holding a lock needed by higher-priority threads
- Convoying
  - Thread holding a lock is preempted, runs out of scheduling quantum, page faults, etc. while holding a lock needed by other threads
- Deadlock
  - Processes attempt to lock the same set of objects in different orders, cyclic dependence

Portland State University – ECE 588/688 – Fall 2010

2

Deadlock Example

- Concurrent Bank transfers

Transfer 1 (from A to B)

Lock(A)  
Load(A)  
Lock(B)→locks A, fails to lock B  
Load(B)  
Calculate new value for A  
Calculate new value for B  
Write(A)  
Unlock(A)  
Write(B)  
Unlock(B)

Transfer 2 (from B to A)

Lock(B)  
Load(B)  
Lock(A) →locks B, fails to lock A  
Load(A)  
Calculate new value for B  
Calculate new value for A  
Write(B)  
Unlock(B)  
Write(A)  
Unlock(A)

Portland State University – ECE 588/688 – Fall 2010

3

Issues with Lock Synchronization (Cont.)

- Livelock
  - Threads that need a lock are starved, unable to acquire it because other threads claim it before they get a chance
- False inter-thread dependencies
  - Conservative programming style can lead to thread serialization, even if it is not really needed
- Performance problems
  - Higher performance requires more fine-grain locking
  - Can lead to more overhead and more false dependencies

Portland State University – ECE 588/688 – Fall 2010

4

Back to Deadlock Example

Transfer 1 (from A to B)

Begin Transaction  
Load(A)  
Load(B)  
Calculate new value for A  
Calculate new value for B  
Write(A)  
Write(B)  
End Transaction

Transfer 2 (from B to A)

Begin Transaction  
Load(B)  
Load(A)  
Calculate new value for B  
Calculate new value for A  
Write(B)  
Write(A)  
End Transaction

- When there are no conflicts, both transactions complete successfully
- When there is a conflict (above), one transaction commits and the other aborts

Portland State University – ECE 588/688 – Fall 2010

5

Solution: Lock-Free Synchronization using Transactional Memory

- Transactional Memory
  - Allows programmers to define customized Read-Modify-Write operations that apply to multiple words of memory
  - Implemented by extending cache coherence protocols
- A *Transaction* is a finite sequence of instructions in a single thread that satisfies two conditions
  - Serializability
    - Transactions appear to execute serially
    - Instructions of one transaction do not interleave with another's
    - Committed transactions are never observed to execute in different orders by different processors
  - Atomicity: All or nothing
    - Each transaction makes tentative changes to memory
    - When completed, a transaction commits (making changes permanent) or aborts (discarding changes) as a whole

Portland State University – ECE 588/688 – Fall 2010

6

### Related Concept: Database Transactions

- Transactions are a widely used concept in database systems
- A database transaction satisfies the ACID properties:
  - ◆ Atomicity: Transaction is executed as a whole, or no part of it is executed (similar to last slide)
  - ◆ Consistency: If database is in a consistent state before transaction, it should be consistent after transaction
  - ◆ Isolation: Concurrent transactions will not interfere with each other's execution. Intermediate changes by a transaction are not seen outside transaction until transaction is committed
  - ◆ Durability: After commit, a transaction's changes are permanent even when system fails
- When a conflict occurs, some transactions are killed to allow others to commit

### Transactional Memory Concepts

- TM primitives
  - ◆ Load-Transactional (LT): reads value of a shared memory location to a private register
  - ◆ Load-Transactional-Exclusive (LTX): reads value of a shared memory location to a private register with the intent to write
  - ◆ Store-Transactional (ST): Tentatively writes a value from a private register to a shared memory location
- Read and write sets
  - ◆ Read set: locations read by LT
  - ◆ Write set: locations accessed by LTX or ST
  - ◆ Transaction's data set: Union of read and write sets

### Changing A Transaction's State

- COMMIT: Attempt to make transaction's tentative changes permanent
  - ◆ A commit succeeds if no other transaction has updated any location in the transaction's data set, and no other transaction has read any location in a transaction's write set
  - ◆ If commit succeeds, all changes to write set are made visible to other threads
  - ◆ If commit fails, all tentative changes to write set are discarded
- ABORT: Discards all updates to a transaction's write set
- VALIDATE: test current transaction status
  - ◆ Successful validate indicates current transaction hasn't aborted (though it may abort later)
  - ◆ Unsuccessful validate indicates a transaction has aborted, discards the transaction's tentative updates

### Suggested Use for Transactions

- Instead of acquiring/releasing locks around critical section, a thread can:
  - ◆ Use LT or LTX to read from a set of locations
  - ◆ Use VALIDATE to check read values are consistent
  - ◆ Use ST to modify a set of locations
  - ◆ Use COMMIT to make changes permanent
  - ◆ If either VALIDATE or COMMIT fails, ABORT and restart
- Can be implemented in software, but hardware implementation is needed for good performance
- Hardware support implies limited transaction size
  - ◆ May trap to software on overflow

### Hardware Implementation Guidelines

- Non-transactional operations use the same caches, cache controllers, and coherence protocols that they would've used in the absence of TM
- Custom hardware support restricted to L1 caches and instructions that communicate with them
- Committing or aborting a transaction is a local operation to the cache, doesn't require communicating with other threads or writing data back to memory

### Example Implementation

- Extends Write-Once snooping coherence protocol
- Each processor maintains two caches
  - ◆ Regular cache for non-transactional operations (direct-mapped)
  - ◆ Transactional cache for transactional operations (fully associative)
    - Similar to Jouppi's Victim cache
    - Holds all tentative writes without propagating them to other processors or memory unless the transaction commits
- Cache Line States: Paper Tables 1 and 2
  - ◆ XCOMMIT lines contain old data, XABORT lines contain tentatively modified data
  - ◆ On Commit, XCOMMIT entries discarded, XABORT entries change to NORMAL
  - ◆ On Abort, XABORT entries discarded, XCOMMIT entries change to normal
- Bus transactions: Paper Table 3

Example Implementation:  
Processor Actions

- Processor maintains two flags
  - ◆ Transaction active (TACTIVE): Whether a transaction is in progress
  - ◆ Transaction status (TSTATUS): Whether transaction is active or aborted
- Non-transactional operations behave like original coherence protocol
- Transactional operations issued by aborted transaction cause no bus cycles, may return arbitrary values
- VALIDATE inst. returns TSTATUS flag
  - ◆ If false, sets TACTIVE to false and TSTATUS to true
- ABORT inst. sets TSTATUS to true and TACTIVE to false
- COMMIT returns TSTATUS, sets TSTATUS to true and TACTIVE to false

Portland State University – ECE 588/688 – Fall 2010

13

Example Implementation:  
Processor Actions (Cont.)

- LT issued by active transaction
  - ◆ Probe Transactional cache for an XABORT entry and return its value.
  - ◆ If hit to NORMAL entry, it changes to XABORT, and an XCOMMIT entry is allocated
  - ◆ If no NORMAL or XABORT entries exist in transactional cache, issue T\_READ cycle on bus. When it completes successfully, set up one XABORT and one XCOMMIT entry in transactional cache
  - ◆ If T\_READ returns BUSY, abort transaction (TSTATUS ← false, drop all XABORT entries, set XCOMMIT entries to NORMAL)
- LTX issued by active transaction
  - ◆ Uses T\_RFO on miss (instead of T\_READ)
  - ◆ Change cache state to RESERVED if T\_RFO succeeds
- ST issued by active transaction
  - ◆ Similar to LTX except that it updates the XABORT entry's data

Portland State University – ECE 588/688 – Fall 2010

14

Example Implementation:  
Cache Actions

- Both regular and transactional caches snoop bus
  - ◆ Ignore all requests for addresses not in the cache
- Regular cache actions
  - ◆ READ or T\_READ: If state is VALID, return value. If state is RESERVED or DIRTY, return value and reset state to VALID
  - ◆ RFO or T\_RFO: return data and invalidate own line
- Transactional cache actions
  - ◆ Acts like regular cache if TSTATUS is false or a request is non-transactional (READ or RFO), except that it ignores entries with transactional tags other than NORMAL
  - ◆ T\_READ: If state is VALID, return value
  - ◆ All other transactional operations: Return BUSY
- Memory responds to WRITE requests
  - ◆ responds to READ, T\_READ, RFO or T\_RFO when no caches do

Portland State University – ECE 588/688 – Fall 2010

15

Performance Evaluation

- Alternatives
  - ◆ Test-and-test-and-set (TTS)
  - ◆ Spin locks with exponential backoff
  - ◆ MCS software queuing (similar to last class's paper)
  - ◆ Hardware queuing: QOSB
    - Add a processor to hardware queue of waiters for a line
    - Allows processor to spin on locally-cached shadow version of line
    - When line is released by processor at head of queue, it is transferred to next waiting processor in queue
  - ◆ Load\_Linked/Store\_Conditional (LL/SC)
    - Load location first (with intent to store)
    - Store a new value only if no updates have occurred to location since load\_linked
- Performance in Paper figures 4, 5, 6

Portland State University – ECE 588/688 – Fall 2010

16

Implementation Issues

- Some disadvantages of original technique
  - ◆ Uses separate, fully-associative transactional cache
  - ◆ Transactional cache size limits transaction length
  - ◆ Many implementation details not discussed
- Data version management
  - ◆ Need to store old and new data modified by a transaction
  - ◆ If using one cache, need to store new data or old data elsewhere
- LogTM
  - ◆ Stores new values in cache, old value in per-thread log in virtual memory
  - ◆ On commit (common case), log is discarded
  - ◆ On abort, old values restored from log by software
- Will TM be successful in making parallel programming easier?

Portland State University – ECE 588/688 – Fall 2010

17

Reading Assignment

- Arthur Veen, "Dataflow Machine Architecture," ACM Computing Surveys, 1986 (Read sections 1, 2, 3 and skim the rest of the paper)
- Gregory Papadopoulos and David Culler, "Monsoon: An Explicit Token-Store Architecture," ISCA, 1990 (Read)

Portland State University – ECE 588/688 – Fall 2010

18